

HPSC

Result presentation

06.01.2025

Audric Demeure and Luca Santoro

Serial Code

Interpolation methods

Let the grid position given by: $i = \lfloor \frac{x}{\Delta x} \rfloor$, $j = \lfloor \frac{y}{\Delta y} \rfloor$.

Nearest Neighbor

Simply rounds the coordinates pair $(\frac{x}{\Delta x}, \frac{y}{\Delta y})$ to the nearest integer index (i, j) and returns the data value at that single grid point:

$$Q(x, y) = Q_{i,j}$$

Bilinear interpolation

Using the interpolation points: $(x_1, y_1) = (i\Delta x, j\Delta y)$ and $(x_2, y_2) = ((i+1)\Delta x, (j+1)\Delta y)$, the weighted coefficients are: $w_x = \frac{x_2 - x}{\Delta x}$, $w_y = \frac{y_2 - y}{\Delta y}$. The interpolated value is then:

$$Q(x, y) = w_x w_y Q_{11} + (1 - w_x) w_y Q_{21} + w_x (1 - w_y) Q_{12} + (1 - w_x) (1 - w_y) Q_{22}$$

Note: $Q(x, y) = Q_{i,j}$ if (i, j) outside domain.

Analysis of η Evaluation

Initial η Evaluation

$$\eta_{i,j}^{n+1} = \eta_{i,j}^n - \Delta t \cdot h_{i,j} \left(\frac{u_{i+1,j} - u_{i,j}}{\Delta x} + \frac{v_{i,j+1} - v_{i,j}}{\Delta y} \right)$$

Corrected η Evaluation

$$\frac{\eta_{i,j}^{n+1} - \eta_{i,j}^n}{\Delta t} = - \frac{h(x_{u_{i+1,j}})u_{i+1,j} - h(x_{u_{i,j}})u_{i,j}}{\Delta x} - \frac{h(x_{v_{i,j+1}})v_{i,j+1} - h(x_{v_{i,j}})v_{i,j}}{\Delta y}$$

$$\Leftrightarrow \eta_{i,j}^{n+1} = \eta_{i,j}^n - \Delta t \left(\frac{h(x_{u_{i+1,j}})u_{i+1,j} - h(x_{u_{i,j}})u_{i,j}}{\Delta x} + \frac{h(x_{v_{i,j+1}})v_{i,j+1} - h(x_{v_{i,j}})v_{i,j}}{\Delta y} \right)$$

Arithmetic Intensity

In the algorithm, the most computationally intensive routines occur in the double-nested loops of `update_eta` and `update_velocities` that loops over $i = 0, \dots, n_x - 1$ and $j = 0, \dots, n_y - 1$.

update_eta:

```
eta_ij = GET(eta, i, j)
- c1_x * (h_ui_plus_1_j * u_ip1_j - h_ui_j * u_i_j)
- c1_y * (h_vi_j_plus_1 * v_i_jp1 - h_vi_j * v_i_j);
```

Roughly 10 FLOPs per grid cell (i,j).

Arithmetic Intensity (ii)

update_velocities

```
u_ij = (1. - c2) * GET(u, i, j)
      - (c1 / param.dx) * (eta_ij - eta_imj);

v_ij = (1. - c2) * GET(v, i, j)
      - (c1 / param.dy) * (eta_ij - eta_ijn);
```

Roughly 8 FLOPs per cell.

Putting both together, we have:

$(10_{\{\text{update_eta}\}} + 8_{\{\text{update_velocities}\}}) * nx * ny = 18 * (nx * ny)$ FLOPs per timestep.

Over nt timesteps, the total is $\approx 18 * nx * ny * nt$ FLOPs.

Arithmetic Intensity (iii)

Memory Accesses per Kernel

- `update_eta` ~13 reads + 1 write \approx 14 total memory operations
- `update_velocities` ~5 reads + 2 writes \approx 7 total memory operations

So combined we have $7 + 14 \approx 21$ memory access. Then each double is 8 byte so we have 21 memory ops per cell $\times 8 \frac{\text{byte}}{\text{op}} = 168$ bytes.

Arithmetic Intensity With 18 FLOPs per cell and 168 bytes per cell:

$$\frac{18 \text{ FLOPs}}{168 \text{ bytes}} \approx 0.11 \text{ FLOPs/byte}$$

Such a low Arithmetic intensity ≈ 0.1 usually indicates **memory-bound behavior**.

Arithmetic Intensity (iv)

Comparing CPU vs. Memory Time

CPU-limited time: $T_{\text{CPU}} = \frac{\text{FLOPs}}{\text{Peak FLOP Rate}} = \frac{18 \times (N_x \times N_y)}{2.8 \times 10^{12} \text{ FLOPs/s}}$

For $N_x = 100$ and $N_y = 100$ and total FLOPs $\approx 1.8 \times 10^5$ per timestep.

Thus we have : $T_{\text{CPU}} \approx \frac{1.8 \times 10^5}{2.8 \times 10^{12}} = 6.4 \times 10^{-8} \text{ s} \approx 0.06 \mu \text{ s}$ per timestep.

Memory-limited time:

$$T_{\text{memory}} = \frac{\text{bytes accessed}}{\text{Memory Bandwidth}} = \frac{168 \times (n_x \times n_y)}{200 \times 10^9 \text{ bytes/s}}$$

with $n_x=n_y=100$, total byte $\approx 168 * 10^4 = 1.68 * 10^6$ bytes.

$$T_{\text{memory}} \approx \frac{1.68 \times 10^6}{2 \times 10^{11}} = 8.4 \times 10^{-6} \text{ s} = 8.4 \mu \text{ s}$$
 per timestep.

$T_{\text{CPU}} \ll T_{\text{memory}}$ so we can conclude that the code is memory-bounded.

Note: $T(n) = \alpha + \beta n \approx \beta n$ since $\alpha \ll \beta n$ (for large data size) \rightarrow focus on the bandwidth-limited part.

Memory Access Bottleneck

Problem: In the original code, loops were written with indices in the wrong order, causing non-contiguous memory access. This led to excessive cache misses and poor performance.

Solution: By reordering the loops to match the fastest-moving index in the innermost loop, the CPU can fetch data more efficiently from cache.

Indexing Order	Speedup
i, j	1.00
j, i	2,41

Keynote: Contiguous memory access leverages cache lines more effectively, drastically reducing memory latency and increasing overall speed.

Numerical Stability

Theoretical CFL Analysis

The CFL (Courant-Friedrichs-Lewy) condition applied to shallow water equations takes the following form:

$$\lambda \propto c \left(\frac{\Delta t}{\Delta x} \right) \leq 1$$

where:

- $c \approx \sqrt{gh}$ represents the characteristic velocity.

In two dimensions, this condition is modified to:

$$\lambda \approx \sqrt{2}c \left(\frac{\Delta t}{\Delta x} \right) \leq 1$$

For our analysis, we selected a maximum bathymetric value of $h_{\max} \approx 21.875$ m, which allows us to express the timestep condition as:

$$\Delta t \leq \frac{\Delta x}{\sqrt{2}c} \approx 0.048\Delta x$$

Numerical Investigation

We conducted multiple simulations with varying Δt and Δx (resp. Δy) to experimentally determine the relationship between Δx and Δt :




Figure 1: Numerical CFL condition determination.

Findings and Implementation

The results confirm that the CFL relationship is indeed linear between Δx and Δt . However, we observe a difference between the theoretical and numerical relationships, which might be partially explained by the approximation $h \rightarrow h_{\max}$.

For the following analysis, we used the following parameters:

- Spatial steps: $\Delta x = \Delta y = 25$
- A more conservative CFL condition:

$$\lambda = c \left(\frac{\Delta t}{\Delta x} \right) \leq \frac{1}{2}$$

This conservative condition was proposed by A.I. Delis and Th. Katsaounis in their 2005 paper “Numerical solution of the two-dimensional shallow water equations by the application of relaxation methods”.

MPI Parallelization

Principle

The computational domain is divided into $P_x \times P_y$ partitions, where each MPI process operates on a distinct subdomain, reducing computation time by distributing the workload while minimizing inter-process communication through strategic splitting in both x and y directions.

Domain Decomposition: Cartesian Topology

- **MPI_Dims_create:** Determines a suitable 2D factorization of the total number of processes, giving us $\text{dims}[0] = P_x$ and $\text{dims}[1] = P_y$.
- **MPI_Cart_create:** Creates a 2D Cartesian communicator (`cart_comm`).
- **MPI_Cart_coords:** Retrieves each process's (x, y) coordinates in the 2D grid.
- **MPI_Cart_shift:** Identifies neighbors (`LEFT`, `RIGHT`, `UP`, `DOWN`) needed for sending/receiving ghost cells.

These function were used to *minimize inter-process communication*.

- **Our configuration:** Each MPI rank holds local arrays for η , u , v , and the interpolated bathymetry `h_interp`.

Local and Global Indices

- Each rank r computes the `start_i`, `start_j`, `end_i`, and `end_j` of its subdomain.
- The local arrays are sized:

$$u \rightarrow (\text{local_nx} + 1) \times \text{local_ny}, \quad v \rightarrow \text{local_nx} \times (\text{local_ny} + 1),$$


$$\text{h_interp} \rightarrow \text{local_nx} \times \text{local_ny} .$$

- Conversion between *global* (i, j) and *local* $(i_{\text{rank}}, j_{\text{rank}})$ is done by:

$$i_{\text{rank}} = i - \text{start_i}, \quad j_{\text{rank}} = j - \text{start_j} .$$

Ghost Cells and Boundary Exchange

- To update η , u , and v near boundaries, we use neighboring data (*ghost cells*).



- MPI_Isend / MPI_Irecv:** Non-blocking send/receive pairs exchange boundary rows or columns with adjacent ranks.
- After posting sends and receives, calls to **MPI_Waitall** ensure data arrival before the next computation step.
- This approach *hides* some communication time by overlapping with computations.

Updating η and u, v in Parallel

- Each rank performs the same numerical updates on η , u , and v , but restricted to its subdomain.

- **update_eta:**

- Uses local arrays plus ghost cells from neighbors:

$$\eta_{i,j}^{n+1} = \eta_{i,j}^n - \Delta t \left[\frac{\partial(h u)}{\partial x} + \frac{\partial(h v)}{\partial y} \right].$$


- **update_velocities:**

- Uses local η -values to compute velocity changes, e.g.

$$u_{i,j}^{n+1} = (1 - \gamma \Delta t) u_{i,j}^n - \frac{g \Delta t}{\Delta x} (\eta_{i,j} - \eta_{i-1,j}) .$$

Gathering and Output

- **MPI_Gatherv:** Rank 0 collects all subdomains' final η -, u -, and v -fields into global arrays.
- **Displacements and sizes:**
 - We compute each rank's portion size and displacement in the global buffer (using arrays like `recv_size_eta` and `displacements_eta`).



Conclusion for MPI Parallelization

- Create a 2D Cartesian topology to split the domain.
- Each process stores only a local subdomain plus ghost cells.
- Non-blocking communications (`MPI_Isend`, `MPI_Irecv`) exchange the necessary boundary information.
- Local updates of η , u , and v proceed in parallel.
- Rank 0 gathers and outputs the final results.
- This method *minimizes communication* by ensuring only boundary data is exchanged among neighbors.

OMP Parallelization

Serial code to OMP Parallelization code

- We used `#pragma omp parallel for` for our main loops in these function:
 - `interp_bathy()`
 - `update_eta()`
 - `update velocities ()`
 - `boundary conditions ()`(partially)
- As far as we tested, (considering only CPU code), using collapse was slowing down our result so we took the decision to not use them.

GPU Parallelization

GPU Parallelization

- **Data mapping:**

- ▶ Fields like `all_data->eta`, `u`, `v` are transferred from CPU to GPU using

```
#pragma omp target teams distribute parallel for collapse(2) map(...) in routines such as  
update_eta() and update_velocities().
```

- ▶ Example in `update_eta()`:

```
#pragma omp target teams distribute parallel for \  
collapse(2)  
map(tofrom: eta_gpu[0:nx*ny]) map(to: h_interp_gpu[], u_gpu[], v_gpu[])
```

- **Boundary conditions** and **source terms** also use `#pragma omp target` to update velocities or apply external forcing on the device.

GPU Parallelization

- We also define a custom mapper for our `data_t` struct to specify how arrays are moved to/from the GPU.
- During each timestep, certain arrays are copied *back* to the host for output:

```
#pragma omp target update from(*all_data->eta, *all_data->u, *all_data->v)
```

- This transfer ensures the most up-to-date GPU results are written to disk.

Nb of GPU	Speedup
1	1.33

Table 2: Speedup against serial code.

Tracing

Scaling Parameters

Following figures were computed using specific parameters to optimize resource utilization:

For the **strong scaling** $dx = 2$, $dy = 2$, $dt = 0.05$ and $max_t = 600$.

For the **weak scaling**: $dx = 5$, $dy = 5$, $dt = 0.02$ and $max_t = 600$. We halved until we reached:
 $dx = 0.3125$, $dy = 0.3125$, $dt = 0.02$ and $max_t = 600$.

MPI Scaling

Strong Scaling





Figure 4: MPI strong scaling 1 and 2 nodes.

MPI Scaling (ii)

Weak Scaling





Figure 5: MPI weak scaling close and spread.

OpenMP Scaling

Strong Scaling





Figure 6: OpenMP strong scaling close and spread.

OpenMP Scaling (ii)

Weak Scaling





Figure 7: OpenMP weak scaling close and spread.

Hybrid Scaling

Strong Scaling






Figure 8: Hybrid strong scaling.

Hybrid Scaling (ii)

Weak Scaling






Figure 9: Hybrid weak scaling.

Hybrid Tracing

Eta analysis

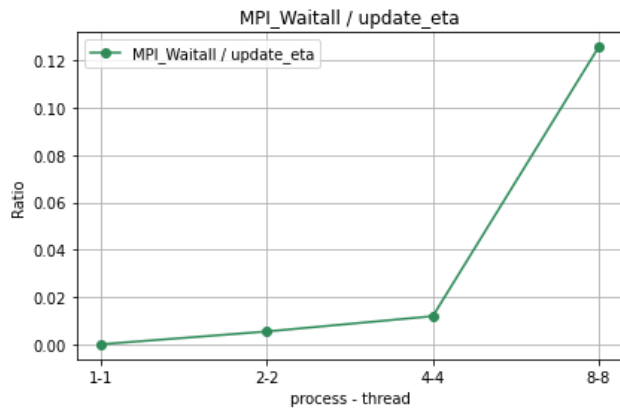
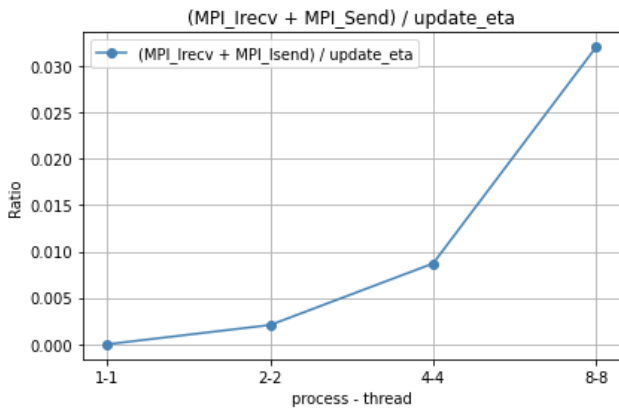


Figure 10: Eta tracing on hybrid code.

Hybrid Tracing (ii)

Velocities analysis

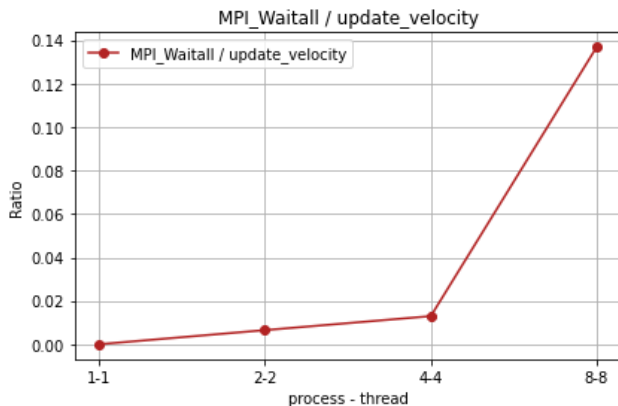
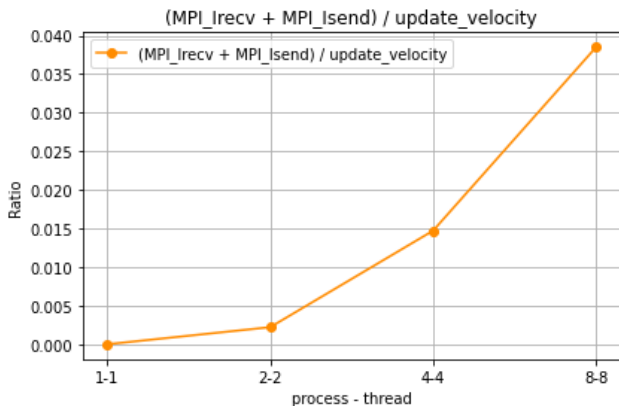


Figure 11: Velocity tracing on hybrid code.

Transparent boundaries

Principle

To simulate the “outer domain”, the transparent boundaries may be implemented using the **Perfectly Matched Layer** (PML) method. Basically, we add a damping term in boundary regions:





Figure 12: PML layer presentation and example on dummy function.

Calculation

The PML is applied to a distance relative to the boundaries:

$$\text{PML}_{\text{width}} \equiv d_{\text{PML}} = \min(0.08N_x, 0.08N_y)$$

The damping amplitude coefficient is assumed dependent of the grid resolution:

$$\sigma_{\max} = 20.0 \left(\frac{25.0}{\Delta x} \right)$$

Cubic damping profile has been used:

$$\sigma = \max(\sigma_x, \sigma_y) \longrightarrow \sigma_{x,y} = \begin{cases} \sigma_{\max} \left(\frac{d_{\text{PML}} - d_{x,y}}{d_{\text{PML}}} \right)^3 & \text{if } d_{x,y} < d_{\text{PML}} \\ 0 & \text{otherwise} \end{cases}$$

Final PML application: $\mathbf{u} \rightarrow \mathbf{u} \times \exp(-\sigma \Delta t)$

Results

Here is depicted the final timestep for the OMP/MPI code with and without the PML applied:




Figure 13: Boundary types comparison with $(\Delta x, \Delta y) = (25m, 25m)$; $\Delta t = 0.05s$ and $t_{\max} = 600s$.
Reflective boundaries on the left and transparent boundaries on the right.

Coriolis Forces

Principle

The previously neglected Coriolis forces are now taken into account:

$$\frac{\partial \mathbf{u}}{\partial t} = -g \nabla \eta - \gamma \mathbf{u} \pm \mathbf{f} \mathbf{u}$$

Where:

- $f = 2\Omega \sin(\varphi) \rightarrow f_{\max} = 2 \times 7.2931 \times 10^{-5}$ (Coriolis coefficient)
- $\Omega = \pi/12$ radians/hour (angular rotation rate of the Earth)

When discretized (example on u):

$$\frac{u_{i,j}^{n+1} - u_{i,j}^n}{\Delta t} = -g \frac{\eta_{i,j}^{n+1} - \eta_{i-1,j}^{n+1}}{\Delta x} - \gamma u_{i,j}^n + f v_{i,j}^n$$

To account for this additional contribution and improve the model, we average actual coupled velocities as:

$$v_{i,j}^n \approx v_{\text{averaged}} = \frac{v_{i-1,j} + v_{i-1,j+1} + v_{i,j} + v_{i,j+1}}{4}$$

Assumptions

Physical Coriolis effects

Realistic Coriolis force ($f = 2 \times 7.2921 \times 10^{-5}$) becomes significant at:

- Spatial scale: $L > 100$ km
- Time scale: $T > 17$ h (inertial period at mid-latitudes)

Our simulation parameters Current numerical setup:

- Domain size: 160×160 points, $dx = dy = 25\text{m} \rightarrow 4\text{km} \times 4\text{km}$ domain
- Time span: $t_{\text{max}} = 600\text{s}$ (10 min), $dt = 0.05\text{s}$

Therefore, with scales 25 times too small spatially and 100 times too short temporally, three options exist:

1. Ignore Coriolis ($f = 0$)
2. Use realistic f (effect negligible)
3. Artificially increase f for testing \rightarrow chosen

This amplification compensates for our reduced scales to observe and validate the implementation:

$$f_{\text{modified}} = 2 \times 7.2921 \times 10^{-2}$$

Results




Figure 14: Water elevation profile at mid-domain.

After 10 minutes of simulation, we can observe modifications in the water elevation pattern: a decrease in water levels on the left side and an increase on the right side. This asymmetric distribution indicates the presence of a rotational motion induced by the Coriolis force.

Erratum

During code review, a critical issue was discovered:

- F_CORIOLIS was not properly defined in the code
- As a result, the Coriolis contribution was exactly zero